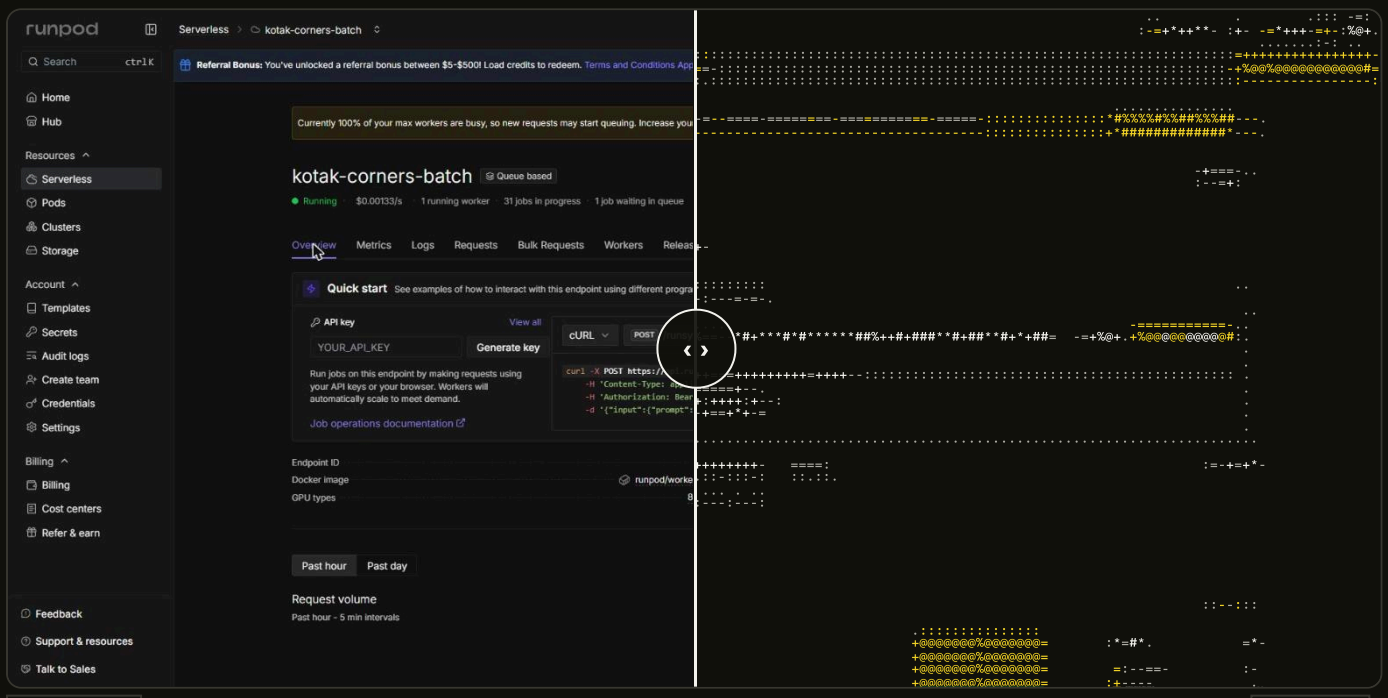
Proof of Use 22 Oct 2026 HTTP errors 0 of 930

RunPod

One serverless H100 run drafted the skill corners for **all 923 O*NET occupations** Kotak checks people against, ahead of time and once.

- kotak-corners-batch
- Qwen3.8-27B-FP8
- 1 × H100 SXM 80 GB



CONSOLE Fig 1, the endpoint overview — the same screen, redrawn in glyphs GLYPHS

01 What ran on RunPod p. 2
02 Handbook checklist p. 2
03 Workload p. 3
04 Example request p. 3
05 Console figures 1–5 p. 4
06 Sources p. 6

923	1,285 s	2.23 M	0
occupations drafted	GPU runtime, smoke + full run	tokens in and out of Qwen3.8-27B	roles failed; 7 answers redrawn

01 What ran on RunPod

Source: RunPod console and `data/generated/role-corners-run.json`

01	**RunPod account**	Team account (AlaskanTuna), funded by the SDC RunPod voucher: US$25 credit redeemed 22 Oct 2026
02	**Product**	Serverless, queue-based endpoint `kotak-corners-batch`
03	**Endpoint ID**	`5yjk19uqmrpbp7`
04	**Endpoint URL**	`https://api.runpod.ai/v2/5yjk19uqmrpbp7` (OpenAI-compatible route `/openai/v1`)
05	**Worker image**	`runpod/worker-v1-vllm:v2.28.0` (vLLM v0.30.0, CUDA 13.0), the stock RunPod vLLM worker
06	**Model**	`Qwen/Qwen3.8-27B-FP8`, open weights under Apache 2.0
07	**GPU**	1 × NVIDIA H100 80GB HBM3 (H100 SXM), data centre AP-IN-1; pools 80 GB PRO and 80 GB; max 1 worker
08	**Runtime**	1,285 s in total: smoke run 502 s (with cold start), full run 783 s
09	**Volume**	930 requests, 0 HTTP errors, 7 answers rejected by the validator and regenerated, 0 roles failed
10	**Tokens**	1,092,522 prompt, 1,137,450 completion
11	**Cost**	About US$1.25: account balance US$25.00 before, US$23.75 after (RunPod console, Billing page, 22 Oct 2026 16:05 MYT)

02 Handbook checklist

Participant Handbook sections 10.7 and 9.4

[x] Endpoint URL, in place of a Pod ID — Fig 1 [1] [2]

10.7: a Serverless team shows its endpoint URL. The endpoint ID and the `api.runpod.ai/v2/5yjk19uqmrpbp7` URL both appear on the endpoint page.

[x] GPU type — Fig 1 [1] · Fig 2 [2]

The endpoint's GPU pools (80 GB Pro, 80 GB), and the H100 SXM the worker actually ran on.

[x] Runtime — Fig 2 [3] · Fig 4 [2]

The worker's uptime mid-run (11 m 16 s) and the endpoint's execution time. The run record gives the exact total: 1,285 s.

[x] The team's own account — Fig 1 [5] · all figures

Every figure is the same signed-in console, with the account's balance and avatar top right.

[x] GPU work runs on RunPod — Section 03, Workload

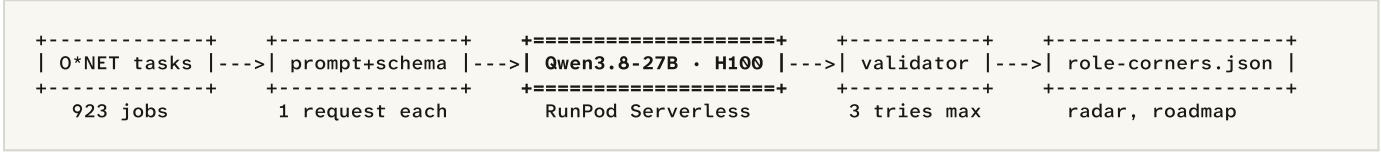
9.4: this batch is Kotak's only GPU work. No model is fine-tuned or self-trained, and the live app calls the Claude and ElevenLabs APIs instead of running a model.

03

Workload

Code: scripts/corners/run-batch.ts
Notes: runpod/README.md

Kotak turns any job into six skill corners and asks twelve yes/no questions about them. RunPod generates that content once, ahead of time, for every one of the 923 O*NET occupations. There is no custom handler: the stock RunPod vLLM worker serves the model. This is Kotak's only GPU work: the live app calls the Claude and ElevenLabs APIs (Technical Declaration, sections 3 and 4).



01	**Input**	The job's title, description and O*NET tasks, grouped by work activity into at most six groups by fixed rules, with no AI.
02	**Ask**	The model groups the tasks into six named corners and writes three can-do statements per corner (basic, working, advanced), with learning hours, SDG targets and search-link resources.
03	**Constrain**	A strict JSON schema goes in `response_format`; vLLM's structured outputs make every token follow it.
04	**Check**	A validator rejects any answer that drops a task, puts a task in two corners or cites a task the job does not have.
05	**Retry**	A rejected answer goes back to the model with its errors. Seven were regenerated; none failed.
06	**Write**	`data/generated/role-corners.json` feeds the check-up questions, the Future Skills Radar and the roadmap; `role-corners-run.json` records every run.

04

Example request

Endpoint settings: runpod/README.md

request.sh — POST /openai/v1/chat/completions

```
curl https://api.runpod.ai/v2/5yjk19uqmrpbp7/openai/v1/chat/completions \
  -H "Authorization: Bearer $RUNPOD_API_KEY" \
  -H "Content-Type: application/json" \
  -d '{
    "model": "Qwen/Qwen3.8-27B-FP8",
    "messages": [{"role": "user", "content": "Name one skill of a graphic designer as JSON."}],
    "temperature": 0,
    "max_tokens": 50,
    "chat_template_kwargs": {"enable_thinking": false},
    "response_format": {"type": "json_schema", "json_schema": {"name": "skill", "schema": {
      "type": "object", "properties": {"skill": {"type": "string"}},
      "required": ["skill"], "additionalProperties": false}}}
  }'
```

rerun.sh — RUNPOD_API_KEY and RUNPOD_ENDPOINT_ID set

```
# smoke run on the three demo roles, then the full run; it skips jobs already done
bun run batch:corners --codes 15-1254.00,15-1255.00,43-3031.00
bun run batch:corners
```

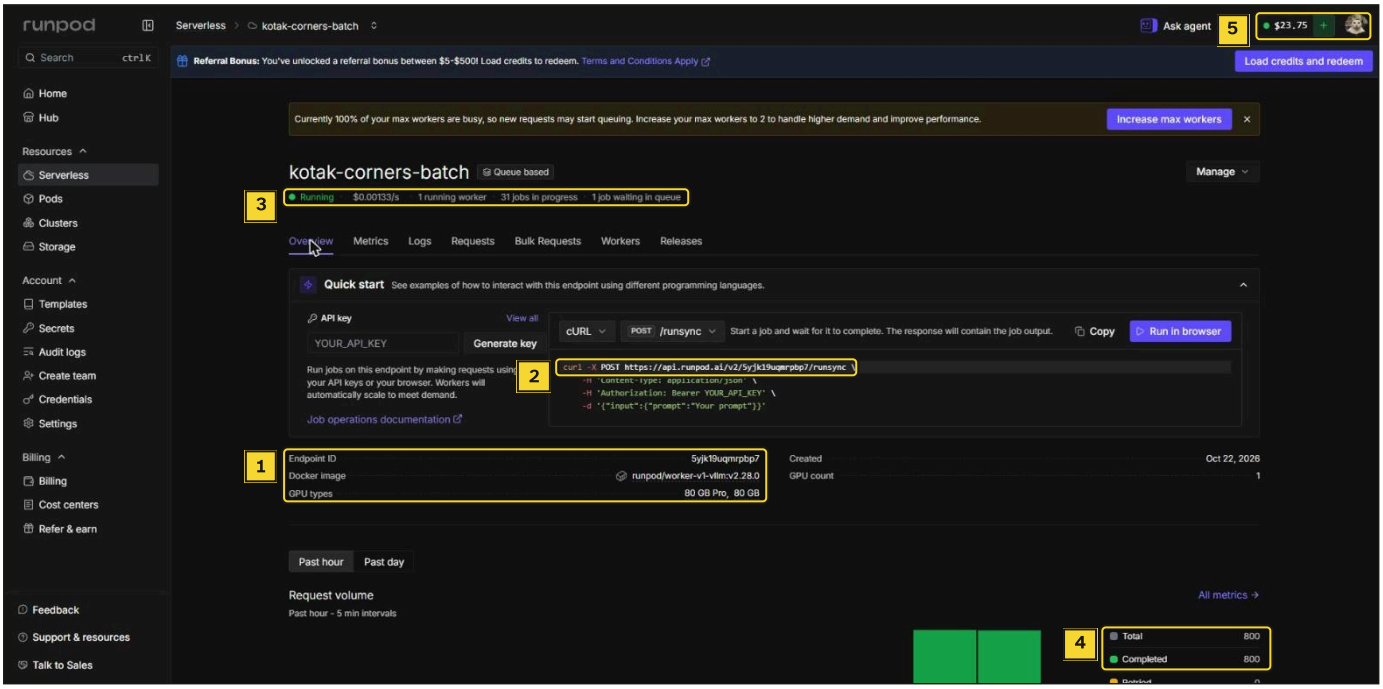
05

Console figures

Captured during the full run,
22 Oct 2026, about 16:00 MYT

FIG 1 Endpoint overview

16:00 MYT



1 Endpoint ID 5yjk19uqmrpbp7, the worker image and the GPU pools: 80 GB Pro, 80 GB.

2 The endpoint URL, in RunPod's own cURL example.

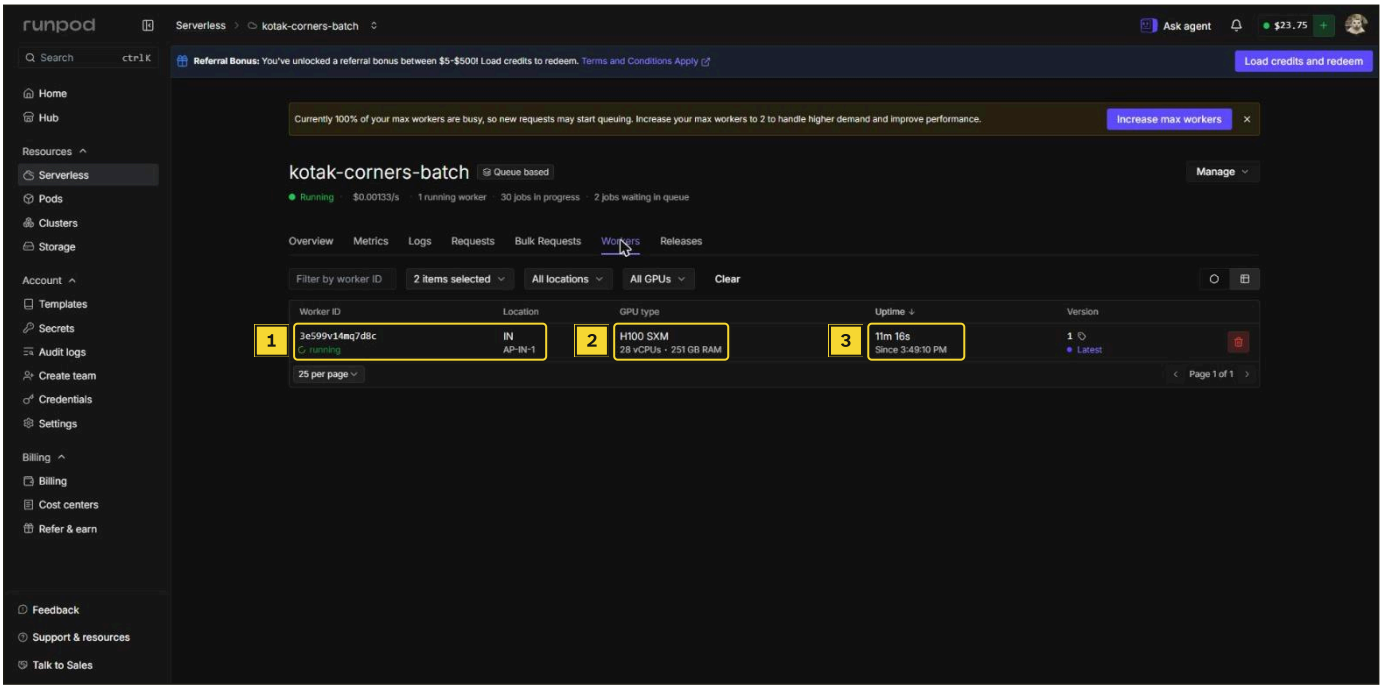
3 Running, with 1 worker and 31 jobs in progress.

4 800 requests completed in the past hour.

5 The account's balance, US$23.75, and avatar.

FIG 2 Workers tab

during the full run



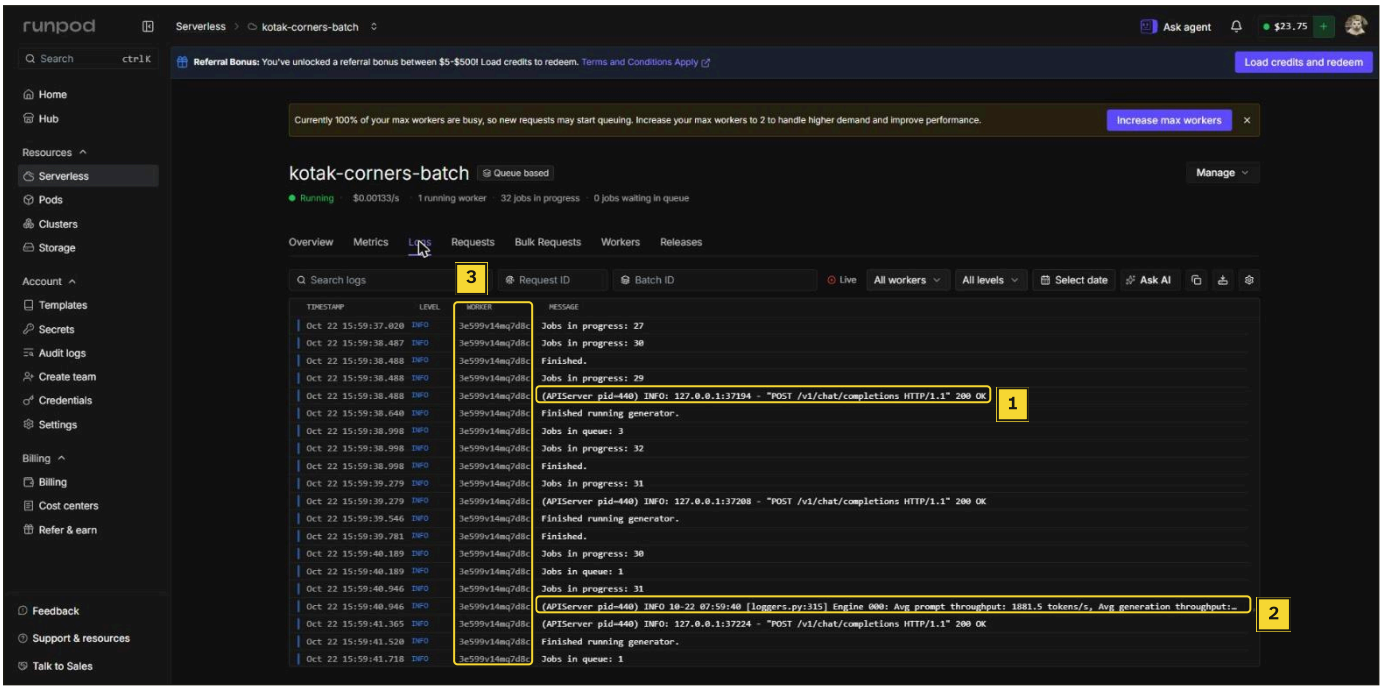
1 Worker 3e599v14mq7d8c, running in AP-IN-1.

2 GPU type: H100 SXM, 28 vCPUs, 251 GB RAM.

3 Uptime 11 m 16 s, since 3:49:10 PM.

FIG 3 Logs tab

15:59 MYT



1 vLLM serving `POST /v1/chat/completions` with 200 OK.

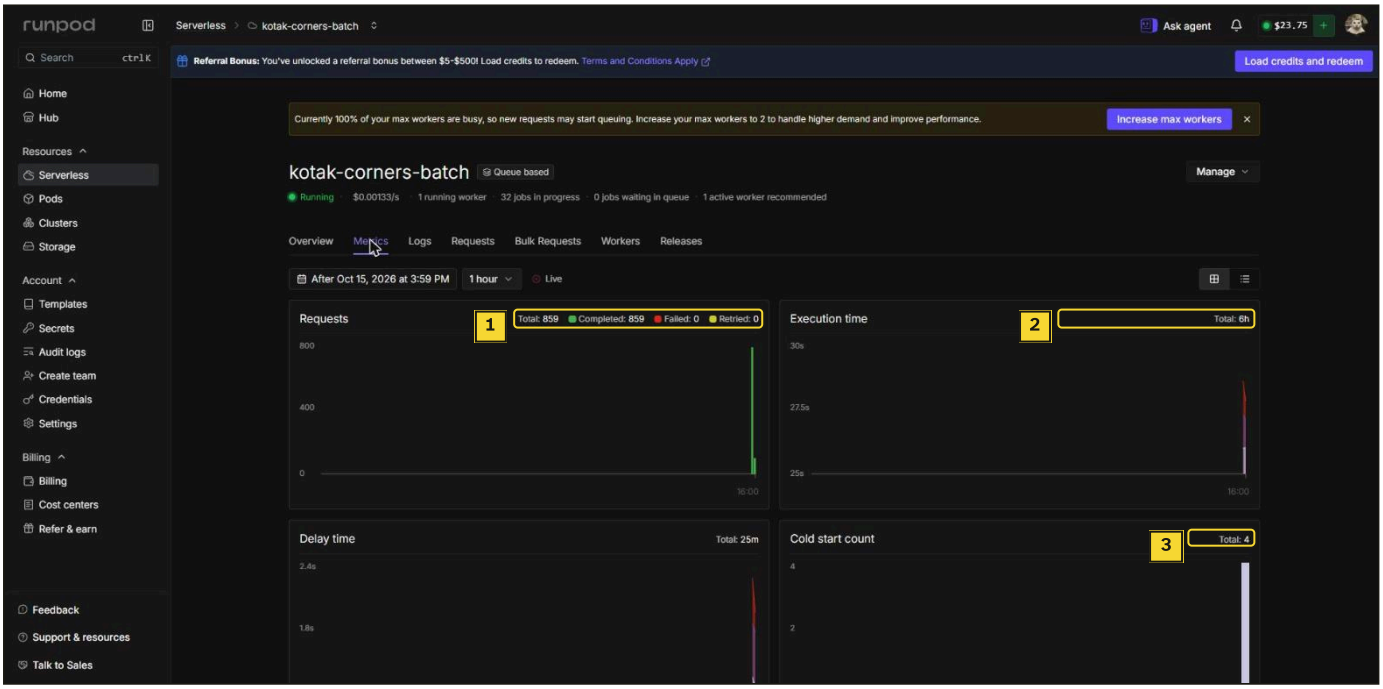
2 The engine's average prompt throughput, 1,881.5 tokens/s.

3 Every line comes from worker 3e599v14mq7d8c.

About 30 jobs are in progress at once: the runner sends 32 requests in parallel.

FIG 4 Metrics tab

at capture time



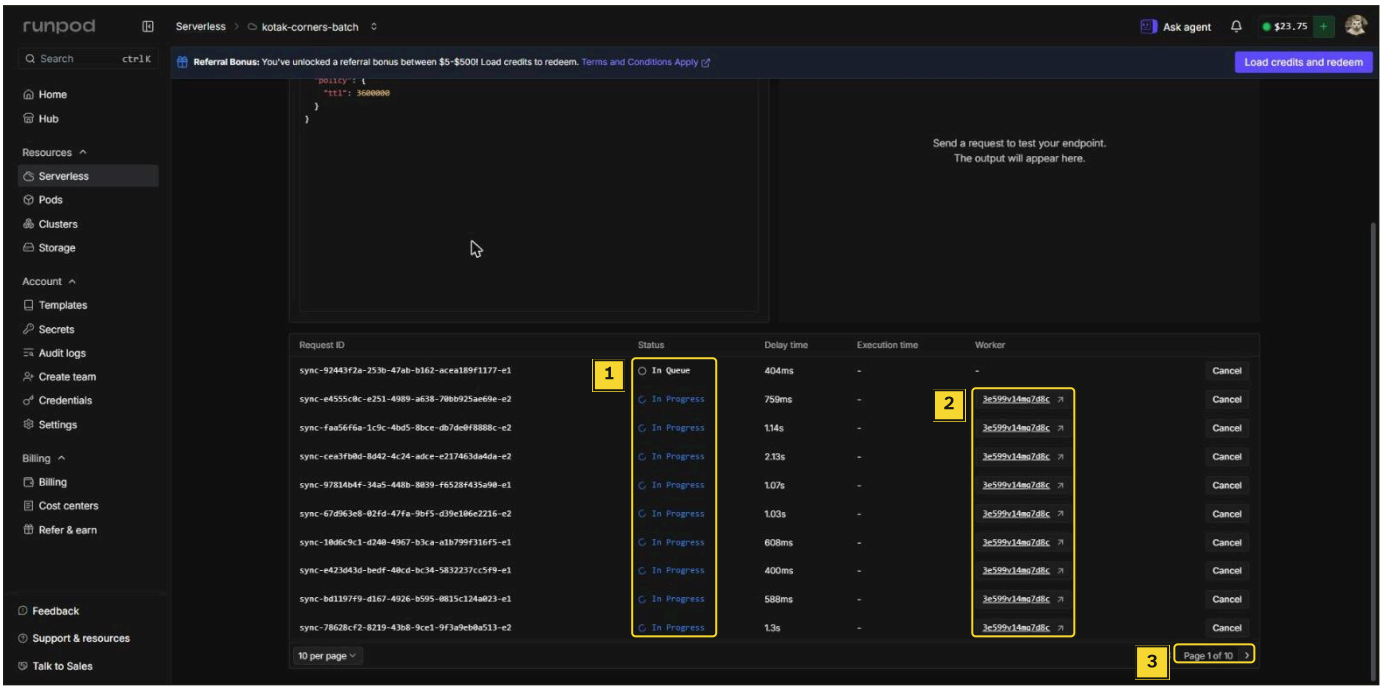
1 859 requests completed, 0 failed, 0 retried.

2 Execution time, 6 h summed over parallel jobs.

3 4 cold starts across both runs.

The filter starts on 15 Oct, so the totals cover both runs; the bars sit at 16:00 on 22 Oct.

FIG 5 **Requests tab** during the full run



1 Batch jobs in progress and in queue.

2 Each one on worker 3e599v14mq7d8c.

3 Page 1 of 10 of the request list.

06 · Sources

Where every number here comes from.

Console	Figures 1–5, 22 Oct 2026
Run record	data/generated/role-corners-run.json
Deploy notes	runpod/README.md
Credit, cost	Handbook 9.1.2; RunPod Billing page
This PDF	scripts/proof-of-use/runpod.html
Repository	github.com/TolongLabs/sdc-hack



