Rotaract Club of Xiamen University Malaysia

Rotary Club of Kuala Lumpur Dynamics

Rotary Club of Tun Razak Exchange

SDC Hackathon

Phase 1 : Preliminary Project Proposal

Hosted by Xiamen University Malaysia (XMUM) AI Club

Submission Deadline: 25 October 2026

Instructions: Complete every field below and submit as a single PDF document named "SDC2026_TeamName_Preliminary.pdf" before the deadline above. Answers must comply with the word limits stated for each open-ended question. Incomplete submissions may not be eligible for review.

Instagram: @sdc_xmumhack ***Email:*** sdc.xmumhack@gmail.com

SDC HACKATHON 2026

Phase 1: Preliminary Project Proposal

Hosted by Xiamen University Malaysia (XMUM) AI Club
Submission Deadline: 25 October 2026

Instructions: Complete every field below and submit as a single PDF document named "SDC2026_TeamName_Preliminary.pdf" before the deadline above. Answers must comply with the word limits stated for each open-ended question. Incomplete submissions may not be eligible for review.

1. Team & Project Overview

Team Name	T010NG
Selected Domain	Domain 1: Aiversecity Future-Ready Workforce
Selected SDG Track (1–17)	SDG 4: Quality Education (target 4.4)
Project / Solution Name	Kotak
Name (Leader)	Hee Zi Jie
Email & Phone (Leader)	heezijie2004@gmail.com, +60125118986
University Affiliation(s)	Asia Pacific University of Technology & Innovation (APU), University of Reading and University of Sheffield
Number of Team Members	5

2. Executive Summary & Problem Statement

2.1 The Problem Context

Describe the specific, real-world challenge your team is tackling. Who is affected by this problem, and why does it matter? Highlight the critical social, environmental, or economic issues involved.

[Maximum 250 words]

AI is changing Malaysian jobs faster than training and skills lists can keep up. In Malaysia, 56% of professionals use AI tools at work, but only 26% have had formal training (Hays Asia Salary Guide 2025, via the Domain 1 brief). In October 2024, the government said about 600,000 workers in 10 sectors needed reskilling within three to five years (Bernama). TalentCorp projects about 620,000 workers, 18% of the workforce in the sectors it studied, will be significantly affected (Impact Study). A World Bank and ISIS Malaysia study rates 28% of the workforce highly exposed to AI (World Bank, 2025). Skills are also under-used: in Q2 2026, 35.2% of employed tertiary-educated workers held semi-skilled or low-skilled jobs (OpenDOSM; hiring-bar evidence).

Our user is a working adult in Malaysia, in any industry, who wonders what AI means for their job. They come with a worry and a few minutes, often on a phone. Some want to keep their job safe, some to move to a role they have in mind, and some to learn where they could go (users).

Today they get static lists. TalentCorp's MyMahir maps roles and skills, but shows no score, time estimate, scenario or SDG tag, and no Malaysian tool we checked joins role data to a scored gap and a timed pathway (prior art). Without dated evidence, a worker cannot tell which skills to learn first, and may spend scarce hours on tasks AI is taking over.

2.2 The Proposed AI Solution

Provide a high-level overview of your proposed application. How does it apply artificial intelligence to solve the problem stated above? Briefly mention the tools or technical skills gained from the SDC Phase 1 workshops (e.g., GenAI APIs, database architecture, cloud hosting) that you intend to leverage.

[Maximum 300 words]

Kotak, Malay for "box", is a web app. Kopi, a cat counsellor, asks 12 yes/no can-do questions about the user's job or a role they want. It draws a Future Skills Radar of the job's 6 core skills, ranks matched roles and builds a dated roadmap in about 3 minutes.

The six must-haves:

1. **Skills profile intake**: a job title in any language, a goal (stay, move or explore) and the 12 questions.
2. **AI-trend signal layer**: dated Anthropic Economic Index snapshots for Malaysia (April and May 2026) and DOSM Job Market Insights, refreshed weekly.

3. **Future-role matching**: ranked roles, each with a fit score, its maths and named gaps.
4. **Reskilling pathway**: ordered steps, each with hours and one free resource, and a finish date set by the user's weekly hours.
5. **SDG mapping**: skills carry SDG targets, jobs show the goals their work serves, and the roadmap highlights steps by goal.
6. **Handling uncertainty**: targets are labelled estimates, Split shows slow and fast AI-uptake futures, and Unknown marks thin data.

Fixed rules, not AI, set every verdict, score, ranking and roadmap order. AI does four jobs:

- Before the demo, a RunPod GPU batch groups each O*NET job's own tasks into 6 skills and drafts the questions. A validator rejects any task outside that job.
- Live, Claude Haiku 5.5 explains each verdict with its dated sources and matches Malay and local job titles.
- Ask Kopi, a chat agent, answers through 6 read-only tools over the same rules, which other assistants reach through an MCP server.
- ElevenLabs voices Kopi, with mute and a text fallback.

From the Phase 1 workshops, we use GenAI APIs (Claude), Postgres (Neon), cloud hosting (Cloudflare Workers and Replit) and RunPod serverless GPUs (workshops; Workshop 1).

3. Technical Architecture & Implementation Plan

3.1 Conceptual Tech Stack

List the technologies, frameworks, APIs, and databases you plan to use to build a functional, deployed prototype. Describe (or attach as an appendix) a simple system block diagram showing how the workshop elements connect into one architecture.

Did your project utilize any official SDC sponsor-provided GenAI API or Cloud Platform? If yes, specify which one and how it is used.

Kotak is a TypeScript monorepo run with Bun. It has two apps, apps/web (the React app) and apps/api (the Hono API), and four packages: shared (the rules and Zod schemas), sdk (the check-up engine), agent (the AI tools) and browser (the research scraper's browser).

- **Frontend**: React and Vite, Tailwind CSS, Recharts for the radar and the roadmap charts, Lucide icons, and pixi-live2d-display for Kopi, built on Live2D's Hijiki sample model. It has phone (390 px) and desktop (1440 px) layouts, follows WCAG 2.2 AA and has a reduced-motion mode (design system).
- **Rule engine in the browser**: the app loads precomputed, dated JSON snapshots and computes levels, targets, verdicts, fit scores and roadmap order in the browser. Switching the snapshot from April to May 2026 recomputes them at once. A guest session in localStorage keeps answers, goals and ticked steps.
- **API**: Hono, on Cloudflare Workers for the main link and on Bun for Replit. One server serves the React build and the API. Its routes match job titles, write cited explanations, proxy text to speech, run Ask Kopi and serve an MCP endpoint. Zod validates every request and response. On Workers, the MCP and chat routes run in a Durable Object.
- **GenAI**: the Claude API with Claude Haiku 5.5. The API caches each answer by a SHA-256 hash of the request and pre-warms the cache for the demo roles. A check rejects any explanation that changes a verdict or lacks a dated citation, and falls back to a fixed template.
- **Agent layer**: @kotak/agent exposes 6 read-only tools over @kotak/sdk: find a role, get the questions, run a check-up, what changed, national skill gaps and list sources. Ask Kopi calls them through Claude, and outside assistants call them through /api/mcp. Every number the agent quotes comes from a tool.
- **Voice**: ElevenLabs. Fixed lines are pre-generated with eleven_v4, and live lines use eleven_v4_turbo. The API caches audio by a hash of the text, and the key stays on the server (TTS notes).
- **Data pipeline**: Bun scripts ingest O*NET 30.1, filter the Economic Index to Malaysia, Singapore and global rows, and extract DOSM Job Market Insights (Q3 2019 to Q2 2026). Every record carries its source, date and layer. A GitHub Actions workflow checks for new releases every Saturday and rebuilds the snapshots.
- **Database**: Neon Postgres 18 (region aws-ap-southeast-1). One table stores anonymous check-ups: only the role code, the 6 skill levels, a sample flag and the date, and totals show only for groups of 10 or more. Other tables hold the research scraper's job listings and its admin accounts.
- **GPU**: a RunPod serverless endpoint runs Qwen/Qwen3.8-27B-FP8 on vLLM for the one-time batch that builds 6 skills and 18 can-do statements for each of 923 jobs.

- **Hosting and CI**: GitHub Actions type-checks, builds and deploys Kotak to Cloudflare Workers on every merge to main, or when a member clicks "Run workflow" (deploy notes). Replit publishes the same main branch from the team's own account. API keys and the database URL live in GitHub, Cloudflare and Replit secrets, never in code.

System block diagram: see Appendix A.

Sponsor platforms: yes, Kotak uses both.

- **Replit**: the team publishes Kotak from main on its own Replit account at the live demo link. The project link, its history and the published app are the proof of use.
- **RunPod**: a serverless GPU endpoint, paid with the team's RunPod credit, ran the batch that built every job's skills and questions: a 3-job smoke test, then 920 jobs in 783 seconds, with none failing. Screenshots of the endpoint, GPU type and runtime are the proof of use (RunPod notes; sponsor rules).

3.2 Design-Thinking & Deployment Strategy

Explain how your team applied design-thinking methodology to ensure this solution is technically viable, user-friendly, and ready for real deployment by Phase 2.

[Maximum 200 words]

- **Empathise**: on 21 October, we named three kinds of worker: stay strong, move to a target role, or ask where to go (users). Research showed Malaysian graduates find work, often below their skill level (hiring-bar evidence).
- **Define**: one user and one moment: a worried worker with a few minutes, who leaves with a plan to start this week.
- **Ideate**: we gave five AI models the brief, as other teams would, and 17 of their 75 ideas served this same worker (field simulation). So Kotak's uniqueness comes from its mechanism: yes/no can-do questions, skills built from the job's own tasks, and verdicts set by visible rules.
- **Prototype and test**: six AI-played experts raised 43 attacks on the first Kotak (roast). We cut the planted-mistake test and moved the demo from a sample card to real data. We set the visual design in Figma before coding.
- **Deploy**: the live demo swaps the Economic Index from April to May 2026, and the radar, ranking and roadmap move. Cached Claude answers and a fixed-template fallback cover API overload. GitHub Actions deploys to Cloudflare Workers on every merge to main, so any member can ship.

4. Preliminary Environmental & SDG Impact Framework

To align with SDC's focus on meaningful global sustainability, establish your project's impact baseline before building the prototype. Use estimated data, published research, or public benchmarks where exact figures are unavailable. Provide the citation of any sources used.

The target, baseline, and KPI you provide here will be used by judges during shortlisting to assess ***Innovation and Creativity*** *as well as* ***Feasibility of build****. Vague or unsupported projections may weaken your case for advancing to Phase 2.*

Metric Category	Your Response
Target SDG Sub-Target(s) *Identify the exact UN SDG sub-target your project addresses* *e.g., Target 11.6 "Reduce the adverse per-capita environmental impact of cities, focusing on air quality and waste management."*	Target 4.4, "By 2030, substantially increase the number of youth and adults who have relevant skills, including technical and vocational skills, for employment, decent jobs and entrepreneurship". It is measured by indicator 4.4.1, "Proportion of youth and adults with information and communications technology (ICT) skills, by type of skill" (UN SDG API, goal 4). The target, indicator and baseline come from the UN SDG API, accessed on 22 October 2026, following our impact notes.
Current Problem Baseline *What is the starting condition of the problem, expressed as a number?* *e.g., "The target infrastructure currently experiences 35% energy inefficiency due to unoptimized scheduling."*	In 2024, 39.8% of Malaysians aged 15 and over had above-basic digital skills in all 5 skill areas, and 19.6% reported writing a computer program in a programming language. Both figures are indicator 4.4.1, series SE_ADT_ACTS, from DOSM (UN SDG API, Malaysia). For AI skills, 56% of Malaysian professionals use AI tools at work, but only 26% have had formal training (Domain 1 brief).
Core Measurable KPI *The specific unit you will track* *e.g., % reduction in power consumption, processing-latency improvement, or kg of resource waste avoided.*	The sourced-roadmap rate. It is the percentage of finished check-ups that reach a roadmap in under 3 minutes, where every verdict, fit score and milestone shows its source and date. A second count tracks uncited recommendations, which must stay at 0 (LLM credibility notes).
Phase 2 Impact Target *Your realistic, percentage-based target by the final prototype presentation* *e.g., "20% improvement in computational resource use."*	100% of 20 seeded test profiles reach a sourced roadmap in under 3 minutes with 0 uncited recommendations (issue #28; LLM credibility notes). On 100% of the hand-checked demo roles, the April-to-May 2026 snapshot swap changes at least one verdict (issue #37).

5. Eligibility & Readiness Checklist

Tick all boxes that apply before submitting.

- ☑ All team members have completed the mandatory GenAI, Database, and Cloud Deployment upskilling workshops.
- ☑ The project concept addresses a verified real-world problem mapped directly to one or more UN SDGs.
- ☑ The proposed solution targets a functional, deployed application rather than a theoretical business case study.
- ☑ All fields in this template have been completed within the word limits stated.

6. IP & Originality Declaration

6.1 Pre-Existing / Prior-Owned Work

Does this submission include any code, design, research, or other work that any team member created or owned before the start of this hackathon?

☐ Yes ☒ No

If yes, briefly describe what was pre-existing and which team member(s) owned it:

Not applicable. Kotak's code, design and research are new work by the team during the hackathon.

6.2 Third-Party Code, Datasets, or Models

Does your project use any third-party code libraries, open-source frameworks, public/private datasets, or pre-trained AI models that your team did not create?

☒ Yes ☐ No

If yes, list them below (open-source use is permitted provided it is declared and properly licensed/attributed):

Name	Type (code/dataset/model)	License
O*NET 30.1 Database	dataset	CC BY 4.0, with O*NET's fixed credit line and a modification notice source
Anthropic Economic Index	dataset	Data CC BY, version not stated; code MIT source
DOSM Job Market Insights	dataset	DOSM data terms: reuse with "Data source: Department of Statistics Malaysia" source
OpenDOSM skill-related underemployment	dataset	CC BY 4.0; one figure cited source
TalentCorp MyCOL and Epoch AI	dataset	Cited as context only, not loaded source
Qwen3.8-27B-FP8, run on RunPod	model	Apache 2.0
Claude Haiku 5.5, through the Claude API	model	Proprietary API service from Anthropic

ElevenLabs text to speech	model	Proprietary API; users keep the rights to the output source
Live2D Hijiki sample model and Cubism Core	model (2D art)	Live2D Free Material License Agreement; Cubism Core under Live2D's own licence
Lucide icons	code	ISC

The Live2D licence says "the Customer may not Redistribute all or part of the Material to third party" (clause 4.1.1). The team keeps the model in the repository and credits Live2D, a risk it accepted on 22 October 2026 (mascot decision).

The code also uses these open-source frameworks and tools: Bun, TypeScript, React, Vite, Tailwind CSS, Recharts, pixi.js, pixi-live2d-display, Hono, Zod, the Anthropic TypeScript SDK, the MCP server SDK, the Neon serverless driver, Cloudflare's Wrangler, Playwright and Clearcote. Every licence is listed in the technical declaration and in NOTICE.md. Kotak runs on hosted services that add no code to the repository: Cloudflare Workers, Neon Postgres, GitHub Actions and Replit.

6.3 Employment-Related IP Conflicts

Is any team member currently bound by an employment, internship, or consulting agreement under which their employer could claim ownership of work they create during this hackathon?

☐ Yes ☒ No

If yes, please describe the nature of the agreement and which member it concerns:

Not applicable. No team member is bound by an employment, internship or consulting agreement that lets an employer claim work made during the hackathon.

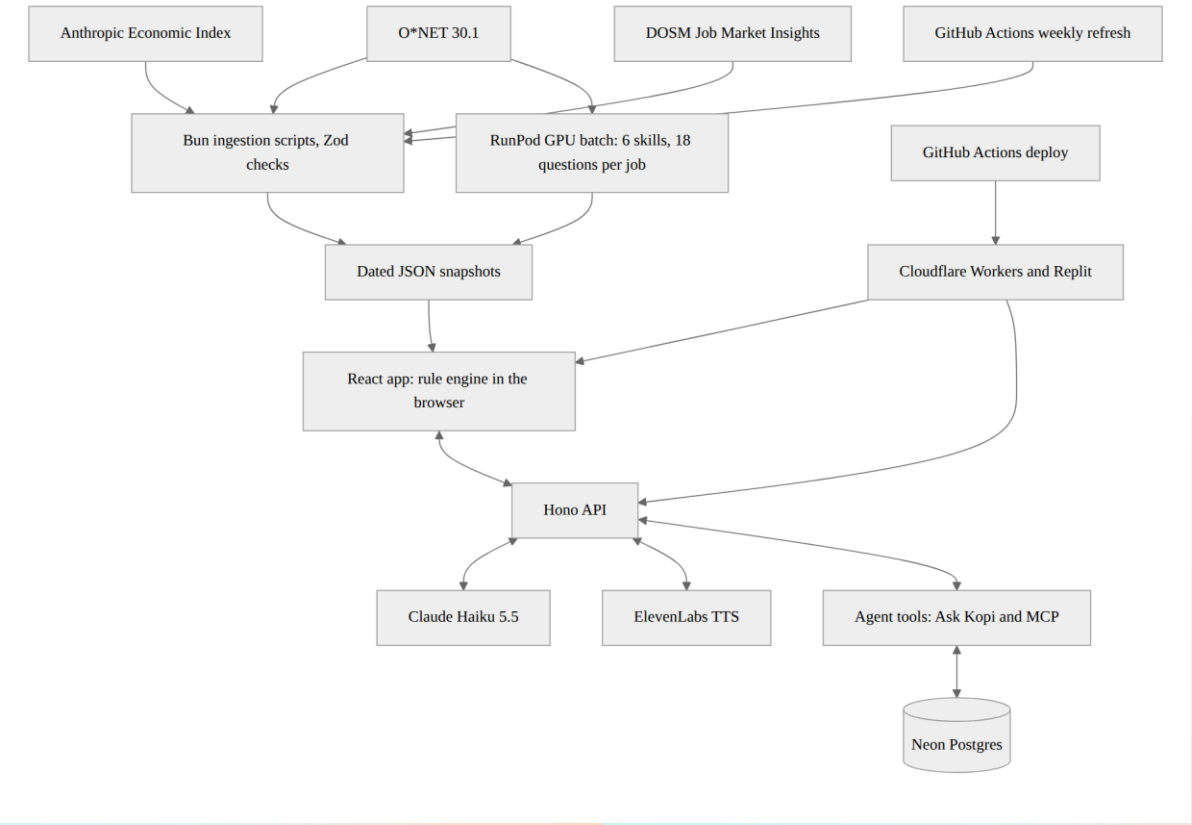
7. Declaration

By submitting this proposal, the team leader confirms that the information provided is accurate to the best of the team’s knowledge and that the project will be original work created during the SDC 2026 hackathon period.



24 October 2026

Appendix A: System Block Diagram



Data flows top to bottom: dated public data and the RunPod batch build the snapshots, the browser computes every verdict, and the API adds Claude, ElevenLabs and the agent tools.